

Can we *learn* the interval-censored Hawkes (MBPP) solver?

A feasibility report on neural surrogates for the covariate-augmented MBPP

Prepared for: Managing Director **Author:** Quantitative Research **Date:** 23 June 2026 **Status:** Experimental feasibility study — synthetic data, CPU-scale

1. Executive summary

We tested whether the classical, slow numerical fitting of our interval-censored Hawkes model (the *Mean Behaviour Poisson Process*, MBPP) can be **replaced or accelerated by learned neural surrogates**. The motivation is speed: calibrating a new sector today repeatedly runs an integral-equation solver inside an optimiser loop; a trained surrogate turns each forward solve into a single network call with exact gradients.

Verdict: yes for the forward problem, with one firm design rule; the inverse problem is harder.

- A **system-conditioned neural operator** learns the MBPP solution map to **3–5 % relative error** and generalises to unseen covariate paths and parameters in **one forward pass (~3 ms)** — a 100×+ speed-up over the iterated solver, with no loss of usable accuracy.
- The single most important finding is a **methodological rule**: the surrogate must be *conditioned on the system* it is solving. Models that try to be one fixed input→output map (a vanilla PINN over parameters, or a Fourier operator fed only the forcing) are **mathematically ill-posed across a family** and plateau at 35–50 % error. Models that ingest the system (DeepONet / physics- informed neural operator) succeed. This is a clean go/no-go we can build on.
- The surrogate is **noise-robust**: it degrades gracefully and only "breaks" (2× the clean error) at ~25 % multiplicative noise on the training targets.
- The **inverse direction** (read the kernel parameters straight off the counts) is fast but only **weakly accurate** at this scale — it recovers the system's overall stability well but not the individual cross-excitation entries, consistent with the known identifiability limits of the model.

Along the way we found and fixed **three code-level bugs** and **one experiment- design flaw** (the parameter grid was sampling physically divergent systems) that were silently blocking the whole programme. The package's analytical core is sound: **53/53 automated tests pass**.

2. Bottom line up front

Task	Approach	Result	Speed	Read
Forward solve, varying covariate path	PINO (physics-informed neural operator, conditioned)	3.1 % median rel. error (held-out)	3.3 ms/solve	✅ Production-promising
Forward solve, fixed system	Fourier Neural Operator	4.9 % median rel. error	<1 ms/solve	✅ Works; noise-robust to ~25 %
Forward solve, <i>unconditioned</i>	Vanilla PINN / FNO on forcing only	35–49 % error	—	❌ Ill-posed — do not use
Inverse: counts → kernel	Amortised CNN	entry corr 0.38; stability corr 0.48	0.2 ms	⚠️ Weak — needs richer signal
Rank which startup is funded	Risk-set softmax ranker	top-5 49 % (oracle 54 %, random 14 %); 88 % of oracle MRR	—	✅ Works — near the achievable ceiling

Relative error = $\| \text{prediction} - \text{exact solver} \| / \| \text{exact solver} \|$ in L2, on held-out test sets. "Exact solver" is our validated numerical MBPP solution.

3. What we set out to test

The MBPP equation we are learning (multivariate, covariate-modulated excitation):

$$\xi(t) = s(t) + \int_0^t K(t,u) \xi(u) du, \quad K_{\{m,j\}}(t,u) = \kappa_{\{m,j\}} \theta e^{\delta \top Z(t)} e^{-\theta(t-u)}.$$

Three programmes, per the project's `PURE_LEARNING.md` specification:

- Forward operator learning** — learn the solution map so a forward solve is a single network call (no integral equation in the optimiser loop).
- Physics-informed solving (PINN / PINO)** — train that map directly from the equation's residual, with no ground-truth solver in the loss, and use it as a fast differentiable surrogate.
- Amortised inverse inference** — learn `counts → (κ, θ, δ, ...)` directly, so calibrating a new sector is a forward pass.

All three are evaluated on the same synthetic "investment family" used in the paper.

4. Results

4.1 The forward operator learns — when conditioned on the system

This is the headline. The same architecture family succeeds or fails depending entirely on **whether the network is told which system it is solving**.


Conditioned (works). The physics-informed neural operator (PINO) ingests the covariate path and parameters and is trained purely from the MBPP residual. On **held-out covariate paths and parameters it never saw**, it reaches:

- **median 3.1 %**, mean 4.4 % relative error;
- **3.3 ms per solve** (vs. an iterated numerical solve in the fitting loop);
- error rises only gently toward the stiff (near-critical) regime.


The Fourier Neural Operator, trained supervised on a **fixed** system, independently confirms the ceiling: **median 4.9 %**, 90th-percentile 8.7 %.

Unconditioned (fails). The vanilla parametric PINN — one network mapping parameters to the solution, with the covariate path held fixed — is **unreliable** (median 35 %, worst cases > 100 %), and training it *longer* makes it worse as a handful of stiff instances dominate the loss. The Fourier operator fed only the forcing $s(t)$ — but with the system matrix varying per sample and **not given to the network** — plateaus at **median 49 %**.

Why. If the system varies but the network can't see it, the same input maps to many different correct outputs — the map is *not a function*. The network can only learn the *average* operator, which is exactly the ~50 % plateau we observe. The fix is architectural: condition on the system (DeepONet/PINO style), or train one operator per fixed system.



Left: error distribution, FNO on a fixed system ($\approx 5\%$) vs. the ill-posed unconditioned family ($\approx 49\%$).
Right: noise robustness — see §4.2.



PINO relative error vs. branching ratio κ on held-out covariate paths: tight and low across the stable range.

4.2 The surrogate is noise-robust

Training the forward operator on **noise-corrupted** targets and testing against the clean solver, the error stays low until the noise is large:


Training-target noise	0 %	10 %	25 %
Clean-test rel. error	5.7 %	7.9 %	16.5 %

The "breaking point" (first level exceeding 2× the noiseless error) is **~25 %** — the operator is not over-fitting to noise and degrades gracefully. This matters because real interval counts are themselves noisy Poisson observations.

4.3 The inverse problem is the hard direction

Reading the $M \times M$ cross-excitation (branching) matrix straight off the interval counts is fast (**0.2 ms/inference**) but only **weakly accurate** at this scale:

- individual matrix-entry correlation **0.38**, mean abs. error 0.13;
- overall **stability (spectral-radius) correlation 0.48** — better than the individual entries.



The pattern is consistent with the model's known identifiability: the **aggregate branching/stability** of a system is far better determined by counts than the **individual cross-excitation entries** are. The practical read: the inverse net is a useful *warm-start / triage* tool, not yet a replacement for the likelihood fit. The forward surrogate (§4.1) is the stronger near-term bet — drop it into the existing fitter to keep the fit's statistical guarantees while removing its cost.

4.4 Application: ranking *which* startup needs funding

A separate, application-facing workstream tests the end-goal directly: a **two-layer marked model** for startup funding. Layer 1 is a weekly **sector** process (a positive-excitation Poisson/Hawkes GLM — *where and when* capital flows); Layer 2 is a dynamic **risk-set ranker** — a conditional softmax over the *live* startups in a sector with a recent-funding "cooldown" covariate — that decides *which* startup receives the round. The question: on synthetic data, can we actually rank the funded startup near the top?

Yes. On a realistic synthetic market (11 sectors, ~35 candidate startups per pick, held-out weeks, 5 seeds), the fitted ranker vs. a random-risk-set baseline — and vs. the **oracle** (the true-parameter ranker, i.e. the best achievable given the data's irreducible stochasticity):

metric	fitted	oracle (ceiling)	random	fitted / oracle
top-1 hit	16.7 %	21.2 %	2.9 %	79 %
top-5 hit	49.2 %	53.7 %	14.3 %	92 %
top-10 hit	69.2 %	71.6 %	28.6 %	97 %
MRR	0.328	0.371	0.118	88 %

The ranker is **3–6× better than chance** and recovers **~80–92 % of the oracle's** top-k / MRR — most of the remaining gap is irreducible noise (each pick is a stochastic draw over ~35 startups), not model error. So the ranking layer works.

One real bug found (and worth fixing). The *forward simulation* (`simulate_marked_paths`) can hang for minutes: `fit_sector_count_model` constrains the lagged sector excitation to be non-negative but **not** its spectral radius, so the fitted sector model can be **supercritical** (we measured $\rho \approx 3.7$). Simulating from it explodes via positive feedback — rates hit the `$\exp(20) \approx 4.85 \times 10^8$` clip, `Poisson( $4.85 \times 10^8$ )` draws hundreds of millions of events, and the per-event inner loop iterates that many times. The *ranking evaluation* uses real observed data (no simulation) and is unaffected — but the end-to-end backtest and its simulation should add a stability projection ($\rho < 1$) on the fitted excitation and a guard on the Poisson event loop. This is the same supercriticality lesson as Issue #1 below.

5. What was broken — and fixed — to get here

These experiments did not run at all on arrival. The blockers and fixes:

#	Issue	Impact	Fix
1	Mis-scoped parameter grid. Excitation is modulated by <code>$\exp(\delta \cdot Z)$</code> , so the <i>effective</i> branching ratio is <code>$\kappa \cdot \exp(\delta \cdot Z)$</code> . The shipped grid let this exceed 1 for ~26 % of samples — physically divergent systems whose target intensity blows up (to $\sim 10^{84}$).	Poisoned training targets and made the "stiff regime" look unlearnable.	Added a stability constraint to all samplers <code>$(\kappa \cdot \exp(\delta \cdot \max Z)) < 1$</code> .
2	JAX bug: a random-seed call overflowed 32-bit integers in modern JAX.	The neural operator (PINO) crashed on launch .	One-line fix.
3	TensorFlow bug: layer weight creation used an outdated Keras API.	Every TF model (FNO, state-space, inverse net) crashed.	Updated to the Keras-3 signature.
4	No data sanitation. ~0.1 % of synthetic systems diverge numerically.	Corrupted the supervised datasets (NaNs) and crashed the inverse-data generator.	Filter non-finite/extreme samples before training.
5	(Prior session) NumPy-2 bug in the optimiser.	Broke every classical fit.	One-line fix.

#	Issue	Impact	Fix
6	Supercritical sector simulator (§4.4). The fitted weekly sector model is non-negative but not spectral-radius-constrained ($\rho \approx 3.7$); simulating from it explodes (rates $\rightarrow \exp(20)$ clip $\rightarrow \text{Poisson}(4.85 \times 10^8) \rightarrow \sim 485$ M-iteration inner loop).	The end-to-end backtest hangs for minutes ; ranking eval (real data) unaffected.	<i>Open</i> : project $p < 1$ on the fitted excitation + cap the Poisson event loop.

After these fixes the analytical core is verified end-to-end: the physics-informed objective is correct to machine precision (residual $\approx 1.4 \times 10^{-14}$) and the full automated suite is green (**53/53**).

6. Recommendations

1. **Adopt the conditioned forward surrogate.** Make the PINO/DeepONet operator the default fast solver and drop it into the interval-censored fit (it is differentiable, so the fit gets exact gradients with no finite differences and no Python solver loop). Expected: large wall-clock reduction at $\leq 5\%$ solve error, while keeping the likelihood fit's statistical properties.
2. **Bake in the stability constraint** everywhere systems are sampled or fit — it is a physical requirement, not a tuning choice.
3. **Treat the amortised inverse as triage, not truth** for now. Use it to warm-start and to estimate system stability; do not yet trust individual cross-excitation entries from it.
4. **Validate at scale on GPU.** All numbers here are **CPU-scale** (small networks, thousands of samples). The design targets GPU with $\sim 10\times$ the data and training; we expect the conditioned operators to improve further and the inverse net to be re-assessed there before any verdict.
5. **Then test on real data.** Everything to date is synthetic, on the paper's data-generating family. The next gate is a back-test on real interval-count series.

7. Reproducibility

Everything is scripted and the artefacts are committed.

- **Environment:** Python 3.12, NumPy 2.x, JAX 0.10.2 (CPU), TensorFlow 2.21 (CPU).
- **Well-posedness gate (NumPy, runs anywhere):** `PYTHONPATH=. python experiments/learning/run_learning.py --selfcheck`
- **Experiment driver (this report's numbers + figures):** `PYTHONPATH=. python experiments/learning/run_deliverable.py --block P|F|A`
- **Metrics:** `results/learning/{P,F,A}_results.json`
- **Figures:** `results/learning/*.png`
- **Test suite:** `python -m pytest tests/ -q` $\rightarrow$ 53 passed.

Headline numbers (verbatim from the JSON artefacts)

Metric	Value
PINO held-out rel. error (median / mean)	0.031 / 0.044
PINO solve time	3.3 ms
FNO fixed-operator rel. error (median / p90)	0.049 / 0.087
FNO noise breaking point	~25 %
Unconditioned FNO (family) median error	0.492
Vanilla PINN median error	0.345
Amortised inverse: entry corr / stability corr	0.38 / 0.48
Fraction of shipped grid that was supercritical	26 %
Automated tests passing	53 / 53

Caveat: this is a feasibility study on synthetic data at CPU scale. The forward- surrogate result is robust and actionable; the inverse result and all absolute accuracies should be re-confirmed on GPU and on real data before production use.